

Advanced Data Mining and Machine Learning

Text Mining

Academic Year 2022-2023

Alessandro Renda

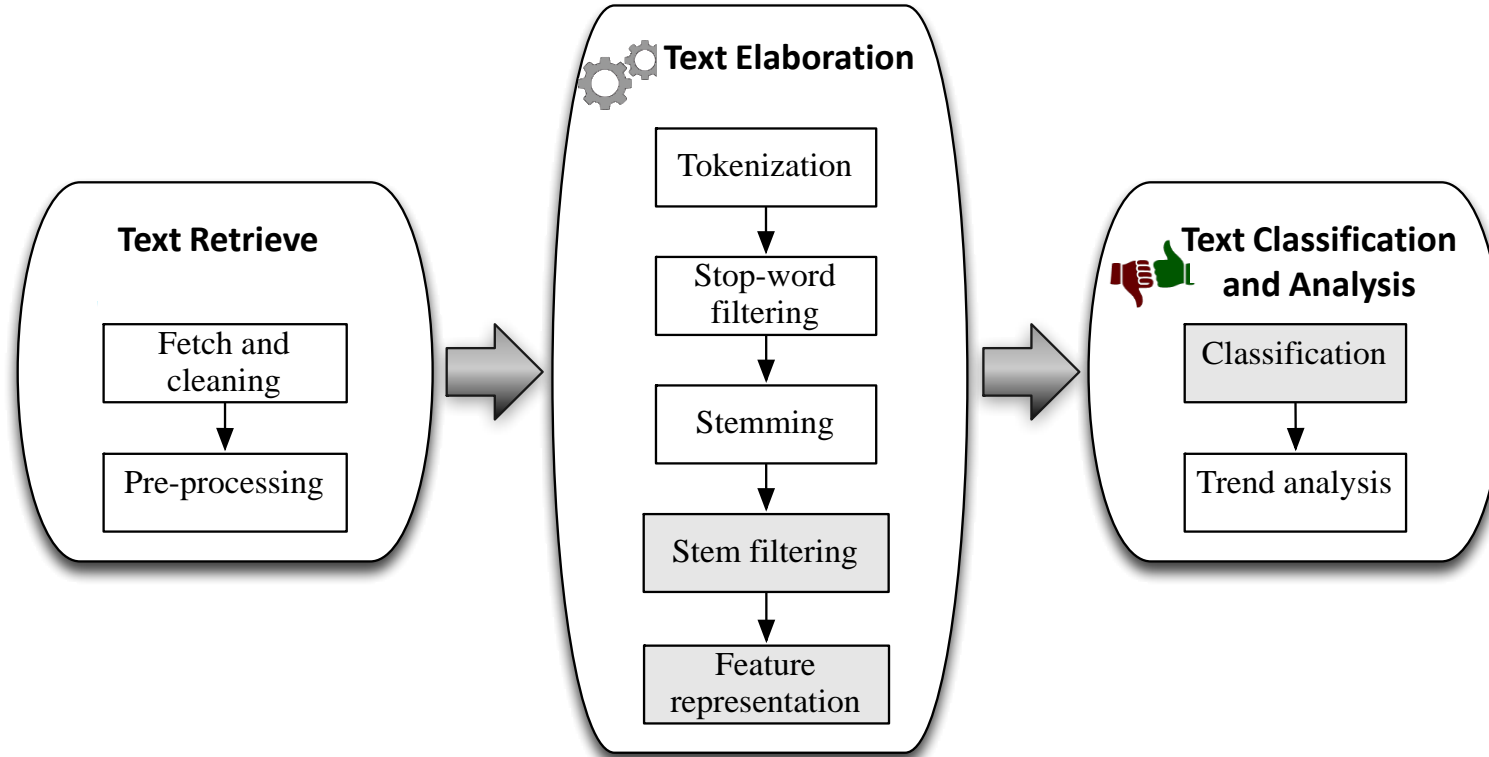
Slides adapted from a collection by professor Pietro Ducange

How can we extract information from TEXTs?

- Text mining refers to the process of ***automatic extraction*** of meaningful information and knowledge from ***unstructured text***;
- ***Text Mining (TM)*** encompasses ***data mining (DM)***, ***machine learning (ML)***, ***statistics***, and ***Natural Language Processing (NLP)***;
- The main difficulty in text mining is caused by the ***vagueness*** of natural language:
 - ***people***, unlike computers, are perfectly able to ***understand*** idioms, grammatical variations, slang expressions, or to contextualize a given word;
 - conversely, ***computers*** have the ability, lacking in humans, to quickly ***process*** large amounts of information.



Text Classification Platform (BOW)



- The platform is completely general. We will describe it relying on a case study of tweets classification

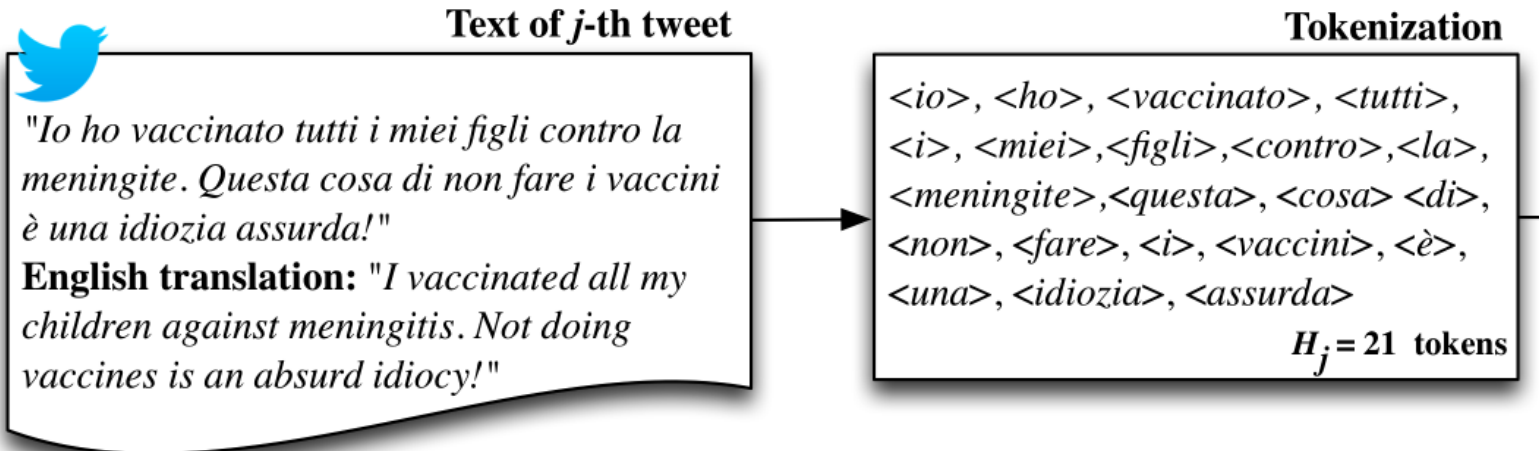
Fetch, Cleaning and Preprocessing

- Streams of text can be fetched from different sources, such as a Micro Blogging Site, based on some ***search criteria*** (e.g., keywords, time or location of posting, hashtags);
- Raw text must be cleaned. For example, we may need to discard:
 - i) ***duplicate*** texts (possibly fetched in different searches)
 - ii) text written in ***languages different*** from the one taken into consideration;
- Text can be preprocessed by applying a ***Regular Expression*** filter, in order to extract ***only the actual text*** and remove all ***useless meta-information***, such as links, hashtags, timestamp, and emoticon.



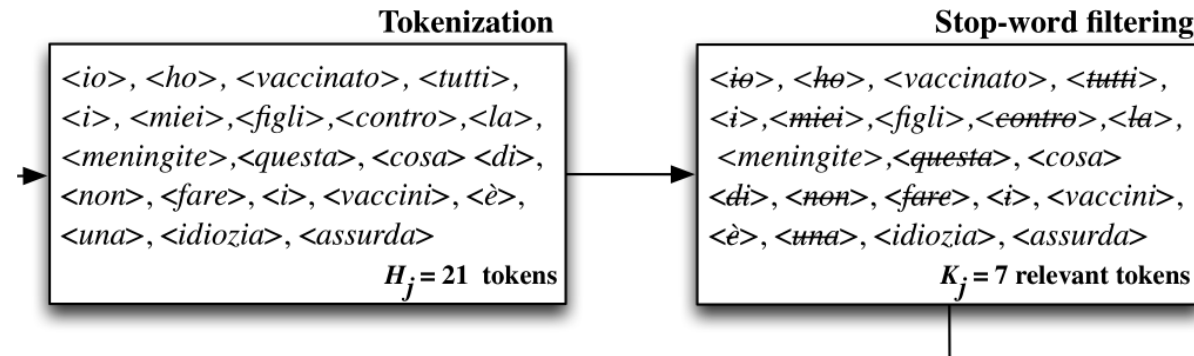
Tokenization

- Tokenization consists in transforming a stream of characters into a stream of ***processing units*** called *tokens*, e.g., words;
- Thus, during this step, after removing punctuation marks, non-text characters and special symbols (e.g., accents, hyphens), each text is represented as a set of words.



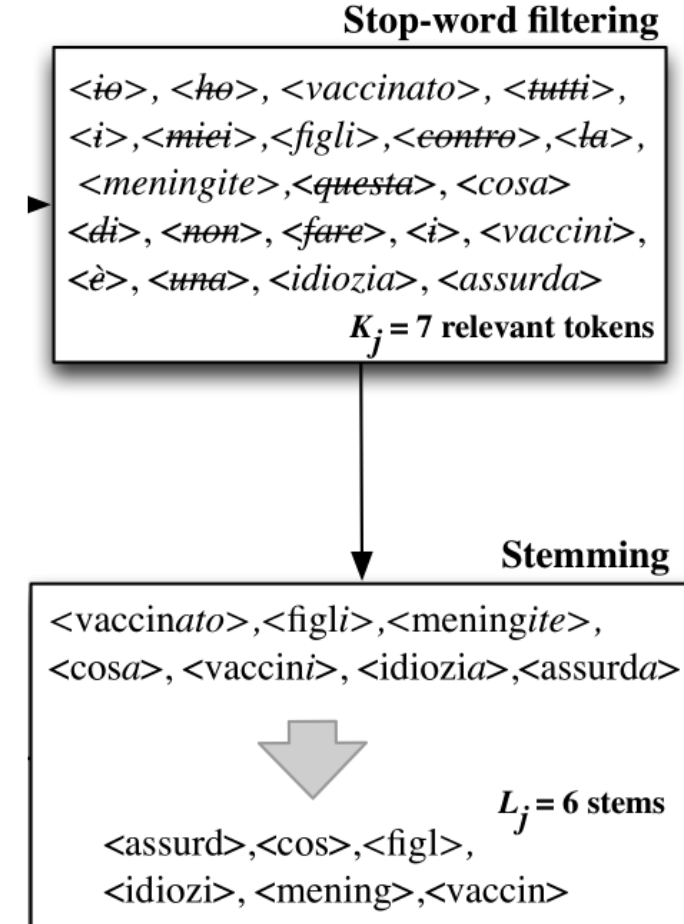
Stop-word Filtering

- This step consists in removing **stop-words**, i.e., words providing little or no useful information to the text analysis, and can hence be considered as **noise**;
- Common stop-words include **articles, conjunctions, prepositions, pronouns**, etc;
- Other stop-words are those typically **appearing very often** in sentences of the considered language (**language-specific stop-words**), or in the particular context analyzed (**domain-specific stop-words**);
- At the end of this step, each text is cleaned from stop-words, and thus reduced to a sequence of **relevant tokens**.



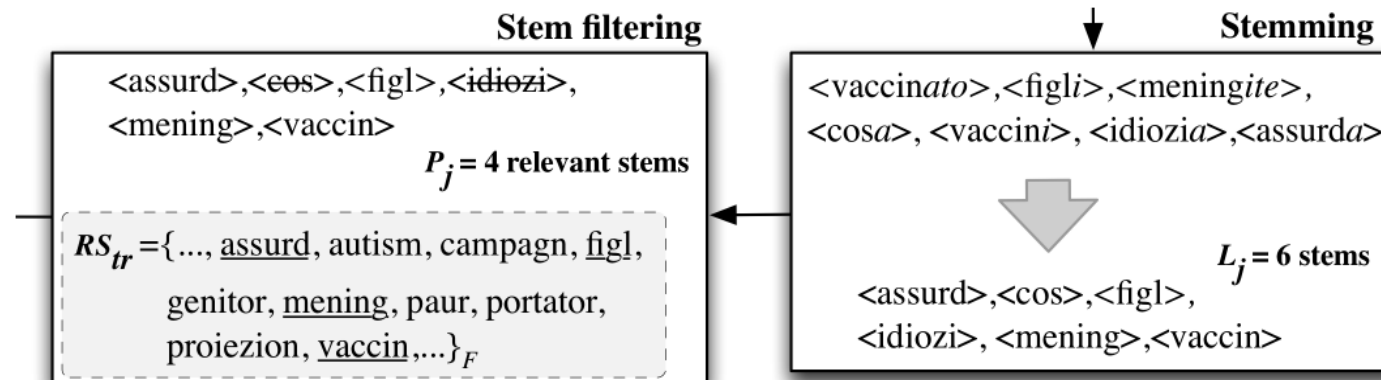
Stemming

- **Stemming** is the process of reducing each token (i.e., word) to its stem or **root form**, by removing its suffix, in order to group words having closely related semantics;
- Hence, at the end of this step each text is represented as a sequence of stems.



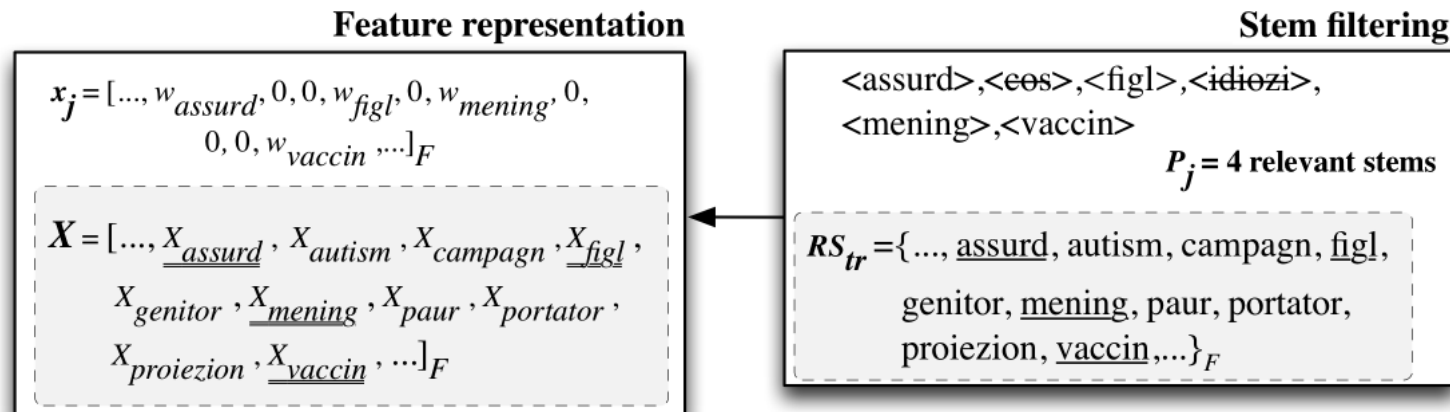
Stem Filtering

- During this stage, the number of stems of each text are reduced, by removing **noisy stems** and maintaining only the most relevant ones;
- Thus, each text is cleaned from stems not belonging to the set of **relevant** stems.
- The set of relevant stem can be provided as a **vocabulary** or identified through a **supervised learning stage**, using the corpus of training documents.

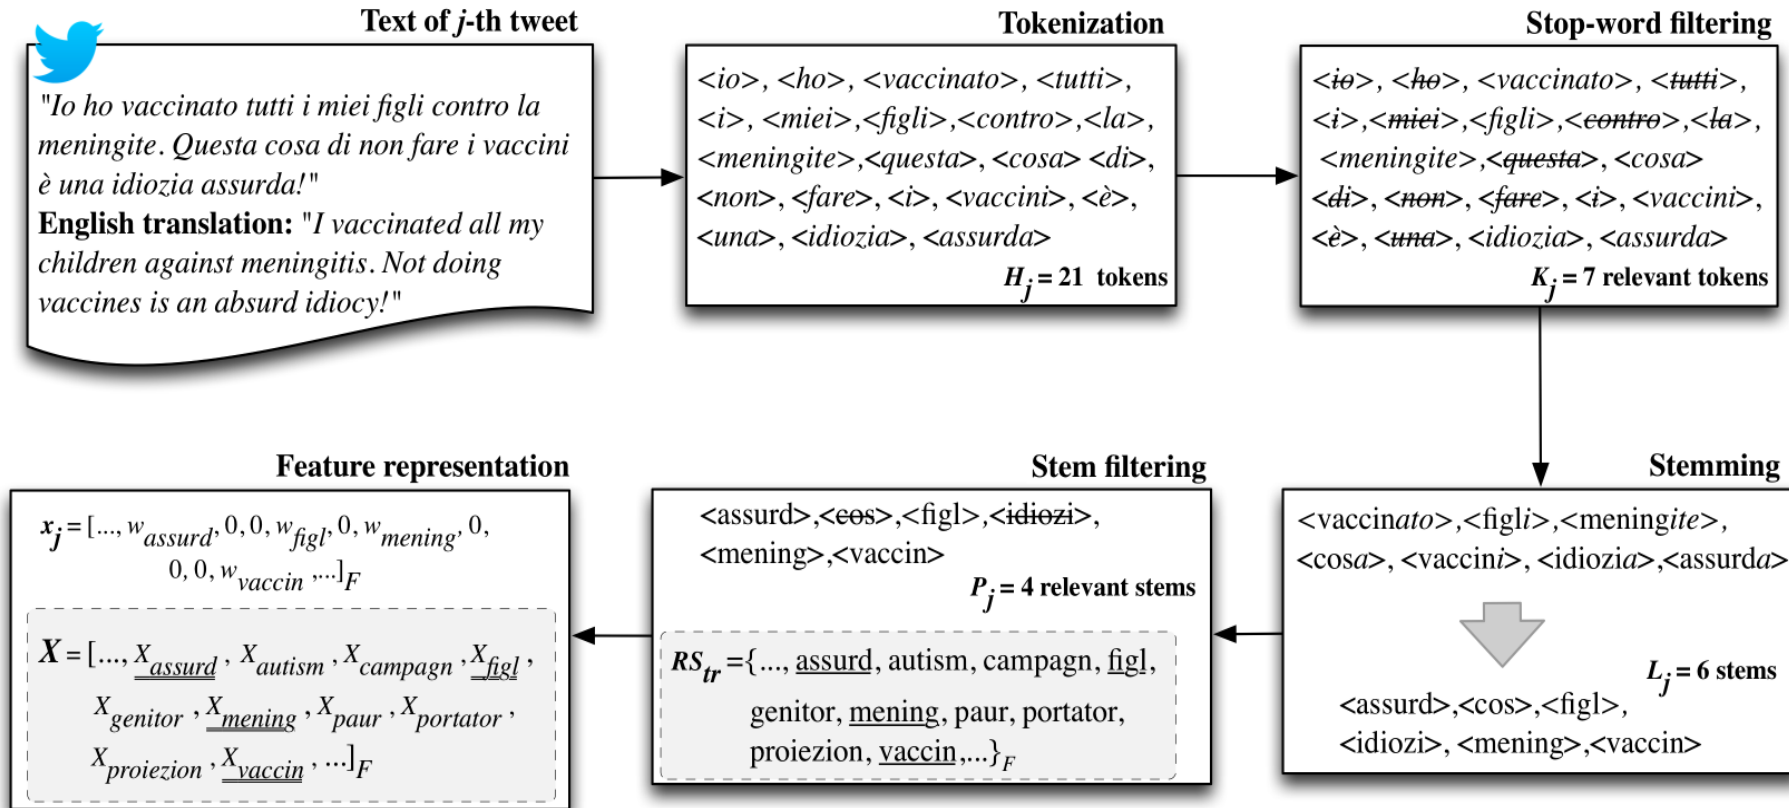


Feature Representation

- **Feature representation** consists in building for each text the corresponding vector of numeric features, i.e., in order to represent all the texts in the same F -dimensional feature space;
- The set of F features corresponds to the set of relevant stems;
- Each text is thus associated with a **vector** of binary or numeric **features**

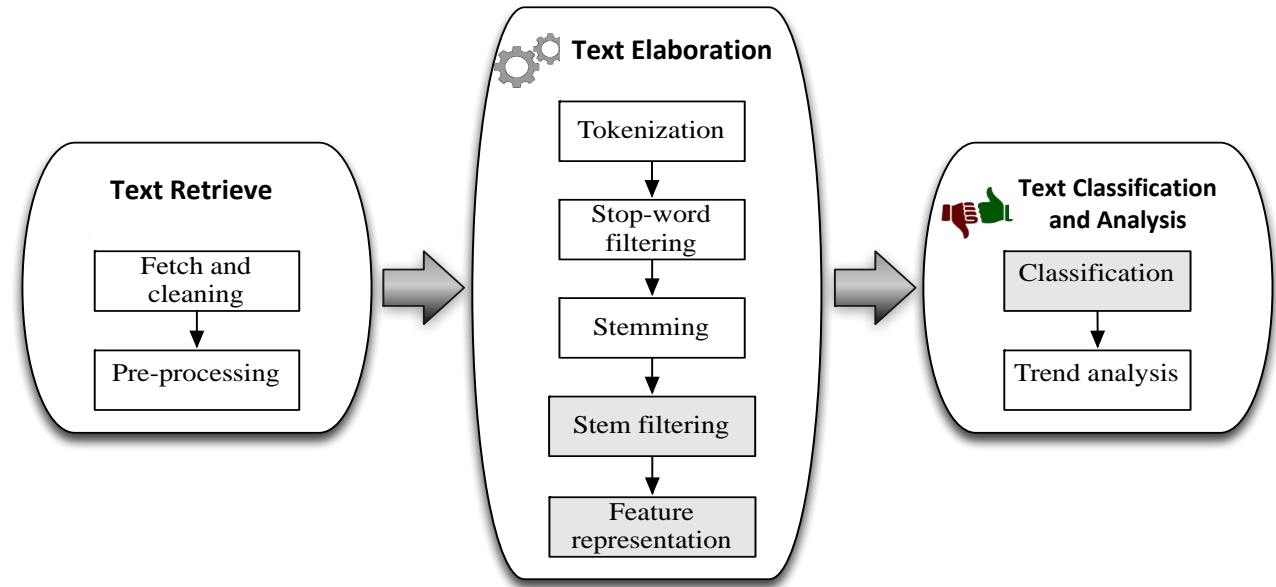


The Entire Text Elaboration Process



The Supervised Learning Stage

- We need to:
 1. Identify the set of **relevant stems**;
 2. Compute the **weights** associated with each of them;
 3. Set the values of the **parameters** of the supervised **classification** model.



- To this aim, we need a collection of N_{tr} labeled texts as training dataset;
- Each text of the training set undergoes the text mining steps: *tokenization*, *stop-word filtering*, and *stemming*;
- The complete set of Q stems is generated **putting together** all the stems extracted from the set of training text after the stemming step.

From text to numbers

- To classify a text, it is necessary to **transform** it into a **numerical vector**, which is then handled by a classification model
- The *Bag-Of-Word* (**BOW**) representation is one of the **simplest** and **most used** technique for text representation:
 - The set of features is composed by the words of the *vocabulary* inferred from the training set
 - Binary, integer or real representation
- Word Embedding (**WE**) methods are also widely used for text representation:
 - Words in a vocabulary are transformed into vectors of continuous real numbers



BOW representation

- The set of F features (vocabulary) corresponds to the set of relevant stems
- Each text is thus associated with a vector of binary or numeric features

Document Vectorization

The quick brown fox jumped over the brown dog

if **binary**:
 "the" <-- 1
if **count**:
 "the" <-- 2

| | | | | | | | | | | | | |
|-----|-----|-------|-------|-----|--------|------|-----|------|------|-----|----------|-------|
| cat | the | quick | brown | fox | jumped | over | dog | bird | flew | ... | kangaroo | house |
| 0 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 0 | 0 | | 0 | 0 |



Dictionary size



TF-IDF: weighting scheme for BOW representation

Let t be a term (e.g., word), d a document (e.g., email), and D the corpus (i.e., collection of documents):

- **Term Frequency (TF)** counts the number of times a term t (word) appears in a document d (i.e., $f_{t,d}$) adjusted by the length of the document (number of all words t' in document d).

$$TF(t, d) = \frac{f_{t,d}}{\sum_{t' \in d} f_{t',d}}$$

- **Inverse Document Frequency (IDF)**: counts the number of documents an individual term t appears over all documents N (inverse fraction of the documents that contain the word, and evaluate the logarithm).

$$IDF(t, D) = \log \frac{N}{|\{d \in D: t \in d\}|}$$

- **Term Frequency - Inverse Document Frequency (IDF)**: product of TF and IDF; weights down common words like "*the*" and gives more weight to rare words like "*software*".

$$TfIdf(t, d, D) = tf(t, d) \cdot idf(t, D)$$



BOW representation drawbacks

- Using BOW representation (with or without TFIDF) results in **large sparse vector** for describing a text.
- This is mainly due to **vast vocabularies** that lead to represent a text by a large vector comprised mostly of zero values.

Document 1

The quick brown fox jumped over the lazy dog's back.

Document 2

Now is the time for all good men to come to the aid of their party.

| Term | Document 1 | Document 2 |
|-------|------------|------------|
| aid | 0 | 1 |
| all | 0 | 1 |
| back | 1 | 0 |
| brown | 1 | 0 |
| come | 0 | 1 |
| dog | 1 | 0 |
| fox | 1 | 0 |
| good | 0 | 1 |
| jump | 1 | 0 |
| lazy | 1 | 0 |
| men | 0 | 1 |
| now | 0 | 1 |
| over | 1 | 0 |
| party | 0 | 1 |
| quick | 1 | 0 |
| their | 0 | 1 |
| time | 0 | 1 |

Image extracted from <https://www.quora.com/What-is-the-bag-of-words-algorithm>



Word Embeddings Representation

Check this example
<https://projector.tensorflow.org/>

- Words are represented by **dense vectors**
- A vector represents the projection of the word into a **continuous vector space**
- **Semantically** similar words have similar vectors
- A word embedding can be learned as part of a **deep learning model**
- A text can be represented using a vector containing of the **average** values of the vectors representing each of its **relevant tokens**.

$$\begin{bmatrix} W_1 \\ W_{11} \\ W_{12} \\ \vdots \\ W_{1n} \end{bmatrix} + \begin{bmatrix} W_2 \\ W_{21} \\ W_{22} \\ \vdots \\ W_{2n} \end{bmatrix} + \dots + \begin{bmatrix} W_n \\ W_{n1} \\ W_{n2} \\ \vdots \\ W_{nn} \end{bmatrix} = \begin{bmatrix} D \\ \frac{W_{11} + W_{21} + \dots + W_{n1}}{n} \\ \vdots \\ \frac{W_{1n} + W_{2n} + \dots + W_{nn}}{n} \end{bmatrix}$$



Word Embedding Learning Methods

- Three popular examples of methods of learning word embeddings from text include:
 - **Word2Vec**: based on Neural Networks (<https://code.google.com/archive/p/word2vec>)
 - **GloVe**: based on matrix factorization (<https://nlp.stanford.edu/projects/glove/>)
 - **FastText**: based on Neural Networks (<https://fasttext.cc/>), created by Facebook
- A number of pre-trained Word Embeddings are available of the websites shown above.
- *Traditional* Word Embeddings (2013) have had a major impact in the field of text mining, but they still have some limits.
- Transformers (2017) and transformer-based language models further improved the state of the art, by enabling the achievement of unprecedented performance in NLP tasks.



Useful References

- Weiss, Sholom M., et al. *Fundamentals of predictive text mining*. Springer, 2015.
- Mikolov, Tomas, et al. "*Distributed representations of words and phrases and their compositionality*." Advances in neural information processing systems. 2013.
- Pennington, Jeffrey, et al. "*Glove: Global vectors for word representation*." Proceedings of the 2014 conference on empirical methods in natural language processing (EMNLP). 2014.
- Bojanowski, Piotr, et al. "*Enriching word vectors with subword information*." arXiv preprint arXiv:1607.04606 (2016).
- D'Andrea, Eleonora, et al. "*Monitoring the public opinion about the vaccination topic from tweets analysis*." Expert Systems with Applications 116 (2019): 209-226.
- Ducange, Pietro, et al. "*An effective Decision Support System for social media listening based on cross-source sentiment analysis models*." Engineering Applications of Artificial Intelligence 78 (2019): 71-85.
- Bechini, Alessio, et al. 2022. *A News-Based Framework for Uncovering and Tracking City Area Profiles: Assessment in Covid-19 Setting*. ACM Trans. Knowl. Discov. Data 16, 6, Article 125, 29 pages.

